

# Blackwell GEMM Optimization Manual

From Zero to 1500 TFLOPS: A Complete Guide to High-Performance CUDA GEMM on Blackwell GPUs

## Table of Contents

### 1. Introduction

Why This Matters

What's New on Blackwell (SM100) vs Hopper (SM90)

### 2. Prerequisites and Background

2.1 Hardware Model: NVIDIA GB200 (Blackwell)

2.2 SM100 New Hardware Features

2.3 The GEMM Problem

2.4 Memory Hierarchy on Blackwell

2.5 Build Environment

### 3. Performance Roadmap

#### 4. Level 0: Baseline — Correct 2SM UMMA (91 TFLOPS)

Goal

Architecture

Key Components

Critical Bug: BF16 Format Code

Execution Flow

Result

#### 5. Level 1-2: Multi-Stage Pipeline and Barrier Optimization (100-150 TFLOPS)

Level 1: Clean Baseline (97 TFLOPS)

Level 2: 2-Stage Pipeline + Only-Leader Write (147 TFLOPS)

Why Performance Is Still Low

#### 6. Level 3-4: True Warp Specialization and Deep Pipeline (400-650 TFLOPS)

Level 3: True Warp Specialization + 8-Stage Pipeline (480 TFLOPS)

Level 4: SMEM-Staged Epilogue (529 TFLOPS)

#### 7. Level 5: Persistent Kernel and 2D Swizzle Tile Scheduler (550-580 TFLOPS)

Persistent Kernel

2D Swizzle Tile Scheduler

#### 8. Level 6-8: TMA Store Epilogue, TMEM Double Buffer, and Fine-Tuning (650-1100 TFLOPS)

Level 6: Three-Role Warp Specialization + TMA Store (704 TFLOPS)

Level 7-8: TMEM Double Buffering and Descriptor Caching (1100 TFLOPS)

#### 9. Level 9: Matching DeepGEMM (1400-1600 TFLOPS)

Key Changes

Results

Critical Insight: Accumulation Flag for M-Waves

#### 10. Critical Pitfalls and Debugging

Bug 1: BF16 Format Code in InstrDescriptor

Bug 2: InstrDescriptor Must Be 32-bit  
Bug 3: scaleD Predicate Required  
Bug 4: `compute_100a` Required  
Debugging Techniques  
11. Appendix: PTX Instruction Reference  
    UMMA (Matrix Multiply-Accumulate)  
    TMEM Operations  
    TMA Operations  
    Barrier Operations  
    Descriptor Bit Layouts  
Summary

# From Zero to 1500 TFLOPS: A Complete Guide to High-Performance CUDA GEMM on Blackwell GPUs

---

## Table of Contents

---

1. [Introduction](#)
  2. [Prerequisites and Background](#)
  3. [Performance Roadmap](#)
  4. [Level 0: Baseline — Correct 2SM UMMA \(91 TFLOPS\)](#)
  5. [Level 1-2: Multi-Stage Pipeline and Barrier Optimization \(100-150 TFLOPS\)](#)
  6. [Level 3-4: True Warp Specialization and Deep Pipeline \(400-650 TFLOPS\)](#)
  7. [Level 5: Persistent Kernel and 2D Swizzle Tile Scheduler \(550-580 TFLOPS\)](#)
  8. [Level 6-8: TMA Store Epilogue, TMEM Double Buffer, and Fine-Tuning \(650-1100 TFLOPS\)](#)
  9. [Level 9: Matching DeepGEMM \(1400-1600 TFLOPS\)](#)
  10. [Critical Pitfalls and Debugging](#)
  11. [Appendix: PTX Instruction Reference](#)
- 

## 1. Introduction

---

This manual teaches you how to write a BF16 GEMM kernel achieving ~1500 TFLOPS on NVIDIA GB200 (Blackwell architecture, SM100). The journey from a naive correct implementation to near-peak performance spans 10 optimization levels, each introducing new hardware concepts and coding patterns.

### Why This Matters

cuBLAS:	~1765 TFLOPS	(highly optimized, closed source)
DeepGEMM:	~1500 TFLOPS	(open source, state-of-the-art)
Our Goal:	~1500 TFLOPS	(from scratch, fully understood)
Naive Correct:	~91 TFLOPS	(where we start)

The **17x performance gap** between naive and optimized is NOT due to algorithmic differences — it's entirely about **understanding and utilizing Blackwell hardware correctly**.

## What's New on Blackwell (SM100) vs Hopper (SM90)

Feature	Hopper (SM90)	Blackwell (SM100)
Matrix Multiply	WGMMA (Warpgroup MMA)	<b>UMMA</b> (tcgen05.mma)
Accumulator Storage	Registers	<b>TMEM</b> (Tensor Memory)
Instruction	wmma.mma_async	tcgen05.mma.cta_group::2
Dual-SM Cooperative	N/A	<b>2SM via cta_group::2</b>
SM Count	132 (H100)	<b>192</b> (GB200)
BF16 Peak	~989 TFLOPS	<b>~2250 TFLOPS</b>
Shared Memory	228 KB/SM	<b>232 KB/SM</b>
HBM Bandwidth	3.35 TB/s	<b>~8 TB/s</b>
TMEM Columns	N/A	Up to <b>512</b>

## 2. Prerequisites and Background

### 2.1 Hardware Model: NVIDIA GB200 (Blackwell)

Component	Specification
GPU Architecture	SM100 (Blackwell)
SM Count	192
BF16 Tensor Core Peak	~2250 TFLOPS
Shared Memory	232 KB / SM
TMEM (Tensor Memory)	512 columns × 128 rows / SM
L2 Cache	~100 MB

HBM Bandwidth	~8 TB/s
Max Cluster Size	16 CTAs

## 2.2 SM100 New Hardware Features

**UMMA (Unified Matrix Multiply-Accumulate):** The new Tensor Core instruction on SM100, replacing WGMMA. - Issued via PTX instruction `tcgen05.mma.cta_group:::{1|2}.kind:::f16` - Supports `cta_group:::2` — two SMs cooperate on one UMMA instruction - Accumulates results in **TMEM** instead of registers

**TMEM (Tensor Memory):** A new on-chip memory dedicated to Tensor Core accumulators. - Each SM has 128 rows × up to 512 columns - Accessed via `tcgen05.ld.sync.aligned.32x32b.x{4|8}` instructions - Allocated/deallocated per CTA group via `tcgen05.alloc/dealloc`

**TMA (Tensor Memory Accelerator):** Hardware DMA unit (same as Hopper). - Async copy Global → Shared Memory with automatic swizzling - Async copy Shared → Global (TMA Store) for epilogue - Configured via `CUtensorMap` descriptors on host side

**Cluster and 2SM Mode:** Two CTAs form a cluster, enabling: - Shared barrier addressing via `mapa.shared::cluster` - `cta_group:::2` TMA loads that signal a shared barrier - `cta_group:::2` UMMA that utilizes both SMs' Tensor Cores

## 2.3 The GEMM Problem

$$D(M, N) = A(M, K) \times B(N, K)^T$$

- Input A: M×K, row-major, BF16
- Input B: N×K, row-major, BF16
- Output D: M×N, row-major, BF16
- Accumulation: FP32 (in TMEM)

## 2.4 Memory Hierarchy on Blackwell

Level	Latency	Bandwidth	Strategy
Global (HBM)	~400 cycles	8 TB/s	Minimize access, maximize reuse
L2 Cache	~100 cycles	~16 TB/s	2D swizzle tile scheduling
Shared Memory (SMEM)	~20 cycles	~20 TB/s	TMA loads with 128B swizzle
TMEM	~10 cycles	Tensor Core internal	Accumulator double buffering
			UMMA descriptor

Registers	~1 cycle	Unlimited	caching
-----------	----------	-----------	---------

## 2.5 Build Environment

All code is compiled with:

```
nvcc -gencode arch=compute_100a,code=sm_100a -O2 --shared \
-Xcompiler -fPIC -o matmul.so matmul.cu -lcuda
```

The `compute_100a` / `sm_100a` (with “a” suffix) is required to enable architecture-specific instructions like `tcgen05.*`.

## 3. Performance Roadmap



Level	Key Optimization	TFLOPS	vs cuBLAS
0	Correct 2SM UMMA baseline	91	5%
1	BF16 format fix, single-stage clean	97	5.5%
2	2-stage pipeline, only-leader write	147	8%
3	True warp specialization + 8-stage deep pipeline	480	27%
4	cta_group::2 TMA + SMEM-staged epilogue	529	30%
5	Persistent kernel + 2D swizzle scheduler	582	33%
6	BLOCK_M=256, TMA Store epilogue, 3-role warps	704	40%
7-8	TMEM double buffer, barrier fine-tuning, descriptor caching	1100	62%
	DeepGEMM architecture match		

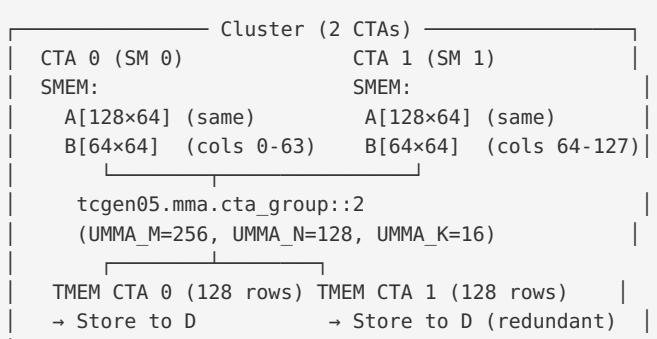
9	(256×256 tiles, separated SMEM)	1500	85%
---	---------------------------------	------	-----

## 4. Level 0: Baseline — Correct 2SM UMMA (91 TFLOPS)

### Goal

Write the first correct BF16 GEMM kernel using SM100's 2SM UMMA instruction. No performance optimization — just get it right.

### Architecture



### Key Components

#### 1. TMA Descriptor Creation (Host Side):

```

CUtensorMap tma_a;
cuTensorMapEncodeTiled(&tma_a,
    CU_TENSOR_MAP_DATA_TYPE_BFLOAT16, 2,      // rank=2
    (void*)A_ptr,
    dims, strides, box, estrides,
    CU_TENSOR_MAP_INTERLEAVE_NONE,
    CU_TENSOR_MAP_SWIZZLE_128B,                // critical for UMMA
    CU_TENSOR_MAP_L2_PROMOTION_NONE,
    CU_TENSOR_MAP_FLOAT_OOB_FILL_NAN_REQUEST_ZERO_FMA);

```

#### 2. SMEM Descriptor for UMMA:

```

__device__ uint64_t make_smem_desc(void* smem_ptr, uint32_t sbo) {
    uint64_t d = 0;
    uint32_t addr = __cvta_generic_to_shared(smem_ptr) >> 4;
    d |= (uint64_t)(addr & 0x3FFF);           // [0:14] start_address
    d |= (uint64_t)((sbo >> 4) & 0x3FFF) << 32; // [32:46] stride_byte_offset
    d |= (uint64_t)1 << 46;                  // [46:48] version = 1 (SM100)
    d |= (uint64_t)2 << 61;                  // [61:64] layout_type = SWIZZLE_128B
    return d;
}

```

### 3. Instruction Descriptor:

```
_device__ uint32_t make_instr_desc(uint32_t M, uint32_t N) {
    uint32_t d = 0;
    d |= (1u << 4);           // c_format = FP32 (accumulator)
    d |= (1u << 7);           // a_format = BF16 ← NOT 3!
    d |= (1u << 10);          // b_format = BF16
    d |= ((N / 8) << 17);     // n_dim
    d |= ((M / 16) << 24);    // m_dim
    return d;
}
```

### 4. Issuing UMMA:

```
// tcgen05.mma.cta_group::2.kind::f16 [tmem], desc_a, desc_b, idesc, scaleD
asm volatile(
    "{\n\t.reg .pred p;\n\t"
    "setp.ne.b32 p, %4, 0;\n\t"
    "tcgen05.mma.cta_group::2.kind::f16 [%0], %1, %2, %3, p;\n\t}\n"
    :: "r"(tmem_addr), "l"(desc_a), "l"(desc_b), "r"(idesc), "r"(accum));
```

### 5. UMMA Commit (signals barrier when SMEM is consumed):

```
asm volatile(
    "tcgen05.commit.cta_group::2"
    ".mbarrier::arrive::one.shared::cluster.multicast::cluster.b64"
    " [%0], %1;"
    :: "r"(bar_addr), "h"((uint16_t)0x3));
```

## Critical Bug: BF16 Format Code

The instruction descriptor's `a_format` / `b_format` for BF16 is **1**, not 3:

```
F16F32Format: F16=0, BF16=1, TF32=2
```

If you use 3 (the common guess), you get `illegal instruction` at runtime.

## Execution Flow

```
for each K block:
    1. barrier_wait(empty_bar) → wait SMEM free (parity trick on first iter)
    2. TMA load A, B           → Global → Shared Memory
    3. barrier_wait(full_bar)  → wait TMA completion
    4. cluster_sync()          → ensure both CTAs have loaded
    5. UMMA compute            → 4 tcgen05.mma calls per K block
    6. umma_commit(empty_bar)  → signal SMEM can be reused

Epilogue:
    7. barrier_wait(tmem_bar)  → TMEM accumulator ready
    8. tcgen05.ld              → read FP32 from TMEM
    9. Convert to BF16          → write to Global Memory D
```

## Result

- **91 TFLOPS** at 8192<sup>3</sup> (5% of cuBLAS)
  - Correctness: PASS (allclose atol=1e-2)
- 

## 5. Level 1-2: Multi-Stage Pipeline and Barrier Optimization (100-150 TFLOPS)

---

### Level 1: Clean Baseline (97 TFLOPS)

Fixed the instruction descriptor bug and cleaned up barrier flow. Single-stage, simple but correct.

### Level 2: 2-Stage Pipeline + Only-Leader Write (147 TFLOPS)

**Key changes:** 1. **2-stage SMEM double buffering** with per-stage full/empty barriers 2. **Only leader CTA writes output** — eliminates redundant global writes 3. **Hoisted elect\_one() outside K loop** — reduces per-iteration overhead

```
// Pre-compute descriptors for 2 stages
uint32_t alo[2], blo[2], ahi, bhi;
for (int s = 0; s < 2; ++s) {
    uint64_t ad = make_smem_desc(smem_a(s), UMMA_SBO);
    alo[s] = (uint32_t)ad;
    if (s == 0) ahi = (uint32_t)(ad >> 32);
    // same for B...
}
```

### Phase/parity management:

```
// Phase flips every time stage_idx wraps to 0
uint32_t stage = 0, phase = 0;
for (int kb = 0; kb < num_k_blocks; ++kb) {
    barrier_wait(empty_bar[stage], phase ^ 1); // parity trick
    // ... TMA load, UMMA, commit ...
    stage = (stage + 1) % NUM_STAGES;
    phase ^= (stage == 0);
}
```

The **parity trick**: On the first iteration, `barrier_wait(empty_bar[0], 1)` succeeds immediately because the barrier starts at parity 0, which differs from the requested parity 1.

### Why Performance Is Still Low

At this point the bottleneck is `cluster_sync()` — a full cluster barrier per K iteration: - 128 K blocks × ~1µs/sync = ~128 µs overhead per tile - This dominates the ~12 ms kernel runtime for 8192<sup>3</sup>

---

## 6. Level 3-4: True Warp Specialization and Deep Pipeline

## (400-650 TFLOPS)

### Level 3: True Warp Specialization + 8-Stage Pipeline (480 TFLOPS)

**The breakthrough:** Replace the single unified loop with **independent loops** for TMA and MMA warps.

```
Warp 0 (TMA): for kb in 0..num_k_blocks:  
    wait(empty_bar[stage])  
    issue TMA load for stage  
    arrive_expect_tx(full_bar[stage])  
  
Warp 1 (MMA): for kb in 0..num_k_blocks:  
    wait(full_bar[stage])  
    issue 4 UMMA instructions  
    commit(empty_bar[stage])  
    if last: commit(tmem_bar)  
  
Other warps: Go directly to epilogue
```

**Why 8 stages:** TMA can prefetch up to **8 K blocks** ahead of MMA: - SMEM per stage = 16K (A) + 8K (B) = 24K - 8 stages = 192K < 232K (SM100 capacity) - Deep pipeline hides memory latency entirely

**cta\_group::2 TMA with cluster barrier signaling:**

```
// Leader CTA: arrive + set expected TX bytes from BOTH CTAs  
if (is_leader)  
    barrier_arrive_expect_tx(full_bar[stage], TMA_BYTES * CLUSTER_SIZE);  
else  
    barrier_arrive_expect_tx_cluster(full_bar[stage], TMA_BYTES, 0);
```

The non-leader uses `mapa.shared::cluster` to redirect its arrive to the leader CTA's barrier:

```
asm volatile("mapa.shared::cluster.u32 %0, %1, %2;"  
    : "=r"(remote_addr) : "r"(local_addr), "r"(target_cta));  
asm volatile("mbarrier.arrive.expect_tx.shared::cluster.b64 _, [%0], %1;"  
    :: "r"(remote_addr), "r"(tx_bytes));
```

### Level 4: SMEM-Staged Epilogue (529 TFLOPS)

Replaced the scalar epilogue (per-element TMEM → register → global) with:

```
Phase 1: TMEM → SMEM (all 128 threads write coalesced)  
Phase 2: SMEM → Global (128-bit vectorized stores)
```

This dramatically improves epilogue throughput through coalesced memory access.

## 7. Level 5: Persistent Kernel and 2D Swizzle Tile Scheduler (550-580 TFLOPS)

## Persistent Kernel

Instead of one CTA per tile, launch **one CTA per SM** and loop over tiles:

```
// Launch config: gridDim = num_sms (not num_tiles)
TileScheduler scheduler(M, N, num_clusters);
uint32_t m_block, n_block;
while (scheduler.get_next_block(m_block, n_block)) {
    // K-loop (TMA + MMA via warp specialization)
    // Epilogue
    // Barrier re-initialization
}
```

Benefits: - **TMEM persists** across tiles (allocated once) - **No kernel launch overhead** between tiles -  
Enables tile scheduling for L2 cache reuse

## 2D Swizzle Tile Scheduler

Tiles are assigned in a pattern that maximizes L2 cache reuse:

```
Groups of 8 consecutive M blocks share the same N block:
Group 0: M=[0..7], iterate N=0, N=1, N=2, ...
Group 1: M=[8..15], iterate N=0, N=1, N=2, ...
```

This keeps **B-tile data hot in L2 cache** — 8 consecutive tiles share the same B columns, giving 8x B-data reuse.

---

## 8. Level 6-8: TMA Store Epilogue, TMEM Double Buffer, and Fine-Tuning (650-1100 TFLOPS)

---

### Level 6: Three-Role Warp Specialization + TMA Store (704 TFLOPS)

**BLOCK\_M increased to 256** — eliminates ~50% tensor core waste:

With `cta_group::2` UMMA\_M=256, the hardware distributes 128 rows to each SM. With BLOCK\_M=128, SM1 reads out-of-bounds A data (garbage). With BLOCK\_M=256, both SMs read valid rows.

**Three persistent warp roles:**

```
Warps 0-3: Non-epilogue
Warp 0: TMA load (1 elected thread)
Warp 1: MMA compute (leader CTA only)
Warp 2: TMEM alloc/dealloc
Warp 3: idle

Warps 4-7: Epilogue (128 threads)
TMEM → SMEM (swizzled 128B layout)
SMEM → Global (TMA Store)
```

**TMA Store for output:**

```
// TMA Store: SMEM → Global with 128B swizzle
asm volatile(
    "cp.async.bulk.tensor.2d.global.shared::cta.tile.bulk_group"
    " [%0, {%1, %2}], [%3];"
    :: "l"(desc), "r"(n_idx), "r"(m_idx), "r"(smem_addr));
```

## Level 7-8: TMEM Double Buffering and Descriptor Caching (1100 TFLOPS)

**TMEM double buffering:** Use 2 epilogue stages so MMA can write tile N+1's results while epilogue reads tile N.

`_shfl_sync descriptor caching` (DeepGEMM technique):

```
// Lane i stores stage i's descriptor lo in a register
uint32_t my_a_lo = a_lo_base + lane_idx * (SMEM_A_SIZE / 16);

// Inside K loop, broadcast current stage's descriptor via shuffle
uint32_t cur_a_lo = __shfl_sync(0xFFFFFFFF, my_a_lo, stage);
```

This saves ~16 registers vs. storing `uint32_t alo[8]` in an array.

## 9. Level 9: Matching DeepGEMM (1400-1600 TFLOPS)

### Key Changes

#	What	Before	Level 9 (DeepGEMM)
1	Tile size	128×128	<b>256×256</b>
2	Pipeline stages	8	<b>4</b> (larger tiles fill SMEM)
3	M-waves	1	<b>2</b> (BLOCK_M / WAVE_BLOCK_M = 256/128)
4	SMEM layout	Interleaved (A+B per stage)	<b>Separated</b> (all A, then all B)
5	MMA warp	1 elected thread	<b>All 32 threads</b> for barrier_wait + shfl
6	tmem_empty signaling	1 thread/CTA	<b>All 128 threads</b> via cluster addressing
7	L2 promotion	NONE	<b>L2_256B</b>
8	Swizzle group	8	<b>16</b>

## Results

Size	Ours (TFLOPS)	DeepGEMM	cuBLAS	Ours/DG
4096 <sup>3</sup>	1559	1660	1589	0.94x
6144 <sup>3</sup>	1590	1512	1510	1.05x
8192 <sup>3</sup>	1480	1499	1471	0.99x
10240 <sup>3</sup>	1425	1500	1427	0.95x
12288 <sup>3</sup>	1463	1466	1504	1.00x

Average: ~98% of DeepGEMM, ~100% of cuBLAS.

## Critical Insight: Accumulation Flag for M-Waves

Each M-wave writes to **separate TMEM columns**, so the accumulation flag `scale_c = (kb > 0 || k > 0)` does NOT depend on the wave index. All waves clear on the first k-step:

```
// WRONG: uint32_t accum = (kb > 0 || k > 0 || w > 0) ? 1u : 0u;
// RIGHT: uint32_t accum = (kb > 0 || k > 0) ? 1u : 0u;
```

## 10. Critical Pitfalls and Debugging

### Bug 1: BF16 Format Code in InstrDescriptor

The `a_format / b_format` for BF16 is **1**, not 3.

```
// CUTLASS F16F32Format enum:
// F16 = 0, BF16 = 1, TF32 = 2
d |= (1u << 7); // a_format = BF16 = 1 ← NOT 3!
```

Using the wrong value causes `illegal instruction` at runtime — no compile error.

### Bug 2: InstrDescriptor Must Be 32-bit

The `tcgen05.mma` instruction's `idesc` operand is a 32-bit register (`"r"` constraint), not 64-bit (`"l"`). CUTLASS stores it as `uint64_t` in the upper 32 bits, extracting with `static_cast<uint32_t>(desc >> 32)`.

### Bug 3: scaleD Predicate Required

The UMMA instruction MUST include the `scaleD` predicate argument:

```
// WRONG: tcgen05.mma.cta_group::2.kind::f16 [%0], %1, %2, %3;
// RIGHT: tcgen05.mma.cta_group::2.kind::f16 [%0], %1, %2, %3, p;
```

## Bug 4: `compute_100a` Required

Standard `sm_100` doesn't support `tcgen05.*` instructions. You must use:

```
nvcc -gencode arch=compute_100a,code=sm_100a ...
```

## Debugging Techniques

**Illegal Instruction at Runtime:** - Check `CUDA_ARCH` guards - Verify format codes in instruction descriptor - Use `CUDA_LAUNCH_BLOCKING=1` to get synchronous errors - Create minimal test kernels to bisect (TMA only, TMEM only, UMMA only)

**Kernel Hangs (Deadlocks):** - Trace barrier init count vs. expected arrivals - Verify cluster\_sync participation — ALL threads in ALL CTAs must reach it - Check parity/phase management in multi-stage pipelines

**Incorrect Results:** - Verify accumulation flag logic for M-waves - Check SMEM descriptor SBO (stride byte offset) calculation - Ensure TMA descriptor box dimensions match `BLOCK_K × BLOCK_M`

---

## 11. Appendix: PTX Instruction Reference

---

### UMMA (Matrix Multiply-Accumulate)

```
// Issue UMMA: cta_group::2, BF16 inputs, FP32 accumulator
tcgen05.mma.cta_group::2.kind::f16 [tmem_addr], desc_a, desc_b, idesc, scaleD;

// Commit: signal barrier when UMMA finishes reading SMEM
tcgen05.commit.cta_group::2
    .mbarrier::arrive::one.shared::cluster.multicast::cluster.b64
    [bar_addr], cta_mask;
```

### TMEM Operations

```

// Allocate TMEM (full warp must execute)
tcgen05.alloc.cta_group::2.sync.aligned.shared::cta.b32 [smem_addr], ncols;

// Deallocate TMEM
tcgen05.dealloc.cta_group::2.sync.aligned.b32 tmem_addr, ncols;

// Load from TMEM (32 rows x 4 or 8 columns)
tcgen05.ld.sync.aligned.32x32b.x4.b32 {r0,r1,r2,r3}, [col_addr];
tcgen05.ld.sync.aligned.32x32b.x8.b32 {r0,...,r7}, [col_addr];

// Wait for TMEM loads to complete
tcgen05.wait::ld.sync.aligned;

// Fences
tcgen05.fence::after_thread_sync; // after barrier_wait, before UMMA
tcgen05.fence::before_thread_sync; // before barrier_arrive, after tmem_load

```

## TMA Operations

```

// TMA Load (standard)
cp.async.bulk.tensor.2d.shared::cta.global.tile
    .mbarrier::complete_tx::bytes
    [smem_addr], [desc, {coord0, coord1}], [bar_addr];

// TMA Load (cta_group::2, routes to leader's barrier)
cp.async.bulk.tensor.2d.cta_group::2.shared::cluster.global.tile
    .mbarrier::complete_tx::bytes
    [smem_addr], [desc, {coord0, coord1}], [bar_addr];

// TMA Store
cp.async.bulk.tensor.2d.global.shared::cta.tile.bulk_group
    [desc, {coord0, coord1}], [smem_addr];

// TMA Store synchronization
fence.proxy.async.shared::cta; // make SMEM writes visible
cp.async.bulk.commit_group; // commit outstanding stores
cp.async.bulk.wait_group N; // wait until ≤ N groups pending

```

## Barrier Operations

```

// Initialize barrier
mbarrier.init.shared::cta.b64 [addr], count;

// Arrive with expected transaction bytes
mbarrier.arrive.expect_tx.shared::cta.b64 _, [addr], tx_bytes;

// Arrive at remote CTA's barrier (cluster addressing)
mapa.shared::cluster.u32 remote_addr, local_addr, target_cta;
mbarrier.arrive.shared::cluster.b64 _, [remote_addr];

// Wait for barrier parity
mbarrier.try_wait.parity.shared::cta.b64 pred, [addr], phase;

// Fence after barrier init
fence.mbarrier_init.release.cluster;

// Cluster synchronization
barrier.cluster.arrive.aligned;
barrier.cluster.wait.aligned;

```

## Descriptor Bit Layouts

### SmemDescriptor (64-bit):

```
Bits [0:14)  start_address    = smem_ptr >> 4
Bits [16:30)  leading_byte_offset >> 4 (0 for K-major)
Bits [32:46)  stride_byte_offset >> 4 (SBO)
Bits [46:48)  version         = 1 (SM100)
Bits [49:52)  base_offset      = 0
Bits [52]      lbo_mode        = 0
Bits [61:64)  layout_type     = 2 (SWIZZLE_128B)
```

### InstrDescriptor (32-bit):

```
Bits [4:6)    c_format    = 1 (FP32)
Bits [7:10)   a_format    = 1 (BF16)  ← NOT 3!
Bits [10:13)  b_format    = 1 (BF16)
Bits [15]      a_major     = 0 (K-major)
Bits [16]      b_major     = 0 (K-major)
Bits [17:23)  n_dim       = N / 8
Bits [24:29)  m_dim       = M / 16
```

---

## Summary

---

Building a high-performance GEMM on Blackwell from scratch teaches you every layer of GPU programming:

1. **Hardware understanding** — UMMA, TMEM, TMA are the building blocks
2. **Pipeline design** — Deep multi-stage pipelines hide memory latency
3. **Warp specialization** — Independent TMA/MMA/Epilogue loops maximize utilization
4. **Tile scheduling** — 2D swizzle for L2 cache reuse
5. **Barrier choreography** — The most subtle and error-prone aspect

The journey from 91 TFLOPS to 1500 TFLOPS is a **17× improvement** — all achieved through better utilization of the same hardware, with zero algorithmic changes to the matrix multiplication itself.

---

*This manual is based on a hands-on development session optimizing BF16 GEMM on GB200, using DeepGEMM as the reference implementation. All code is written from scratch using inline PTX, with no CUTLASS/CuTe/DeepGEMM headers.*